

# Homework 2: Exploratory data analysis and visualization

UIC CS 418, Spring 2022

*According to the **Academic Integrity Policy** of this course, all work submitted for grading must be done individually, unless otherwise specified. While we encourage you to talk to your peers and learn from them, this interaction must be superficial with regards to all work submitted for grading. This means you cannot work in teams, you cannot work side-by-side, you cannot submit someone else's work (partial or complete) as your own. In particular, note that you are guilty of academic dishonesty if you extend or receive any kind of unauthorized assistance. Absolutely no transfer of program code between students is permitted (paper or electronic), and you may not solicit code from family, friends, or online forums. Other examples of academic dishonesty include emailing your program to another student, copying-pasting code from the internet, working in a group on a homework assignment, and allowing a tutor, TA, or another individual to write an answer for you. Academic dishonesty is unacceptable, and penalties range from failure to expulsion from the university; cases are handled via the official student conduct process described at <https://dos.uic.edu/conductforstudents.shtml> (<https://dos.uic.edu/conductforstudents.shtml>).*

This homework is an **individual assignment for all graduate students. Undergraduate students are allowed to work in pairs and submit one homework assignment per pair.** There will be no extra credit given to undergraduate students who choose to work alone. The pairs of students who choose to work together and submit one homework assignment together still need to abide by the Academic Integrity Policy and not share or receive help from others (except each other).

There are three parts in this homework. The first one is a practice introduction to `matplotlib` (5%). The second is a guided exploration of a bikeshare dataset (45%). The third one is a self-guided exploration of a dataset on social media and internet attitudes (50%). You can also earn extra credit of 20%.

## Due Date

This assignment is due at 11:59pm Friday, February 18th.

## What to Submit

You need to complete all code and answer all questions denoted by **Q#** (each one is under a bike image) in this notebook. When you are done, you should export **hw2.ipynb** with your answers as a PDF file, upload the PDF file to *Homework 2 - Written Part* on Gradescope, tagging each question. You need to upload a completed Jupyter notebook (hw2.ipynb file) to *Homework 2 - code* on Gradescope. If one of these two parts (written and code) is missing, you will lose 50%. For undergraduate students who work in a team of two, only one student needs to submit the homework and just tag the other student on Gradescope.

## Autograding

We will not use autograding for this homework assignment.

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
%matplotlib inline
import matplotlib.pyplot as plt
```

## Practice: matplotlib (5%)

`matplotlib` (<http://matplotlib.org/>) is the most widely used plotting library available for Python. It comes with a good amount of out-of-the-box functionality and is highly customizable. Most other plotting libraries in Python provide simpler ways to generate complicated `matplotlib` plots, including `seaborn`, so it's worth learning a bit about `matplotlib` now.

Notice how all of our notebooks have lines that look like:

```
%matplotlib inline
import matplotlib.pyplot as plt
```

The `%matplotlib inline` magic command tells `matplotlib` to render the plots directly onto the notebook (by default it will open a new window with the plot).


Then, the `import` line lets us call `matplotlib` functions using `plt.<func>`

Here's a graph of `cos(x)` from 0 to  $2 * \pi$ .

```
In [2]: # Set up (x, y) pairs from 0 to 2*pi
xs = np.linspace(0, 2 * np.pi, 300)
ys = np.cos(xs)

# plt.plot takes in x-values and y-values and plots them as a line
plt.plot(xs, ys)
```


Out[2]: [



`matplotlib` also conveniently has the ability to plot multiple things on the same plot. Just call `plt.plot` multiple times in the same cell:

```
In [3]: plt.plot(xs, ys)
plt.plot(xs, np.sin(xs))
```

Out[3]: [



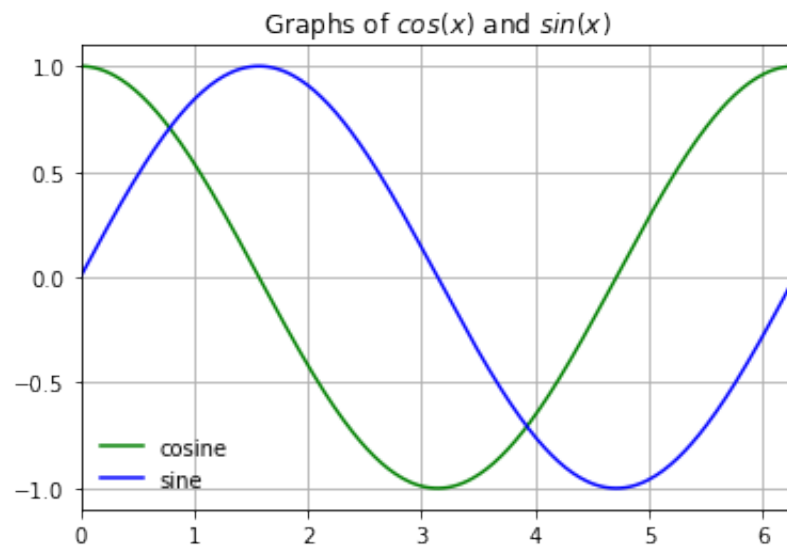
That plot looks pretty nice but isn't presentation-ready. Luckily, `matplotlib` has a wide array of plot customizations.



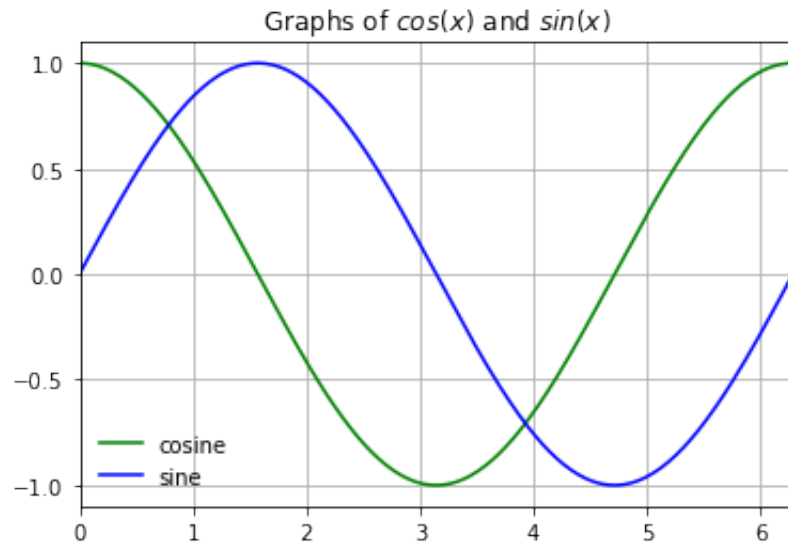
## Q0 (5%):

Skim through the first part of the tutorial at <https://github.com/rougier/matplotlib-tutorial> (<https://github.com/rougier/matplotlib-tutorial>) to create the plot below. There is a lot of extra information there which we suggest you read on your own time. For now, just look for what you need to make the plot.

Specifically, you'll have to change the x and y limits, add a title, and add a legend.



```
In [4]: # Here's the starting code from last time. Edit / Add code to create t
plt.plot(xs, ys, color="green", label="cosine")
plt.plot(xs, np.sin(xs), color="blue", label="sine")
plt.xlim(0,6.3)
plt.yticks(np.arange(-1, 1.1, 0.5))
plt.title("Graphs of  $\cos(x)$  and  $\sin(x)$ ")
plt.legend(frameon=False)
plt.grid()
```



## Part 1: Guided EDA of bikeshare trips (40%)

You will be performing some basic EDA (exploratory data analysis) on bikeshare data in Washington D.C.

The variables in this data frame are defined as:

- instant: record index
- dteday : date
- season : season (1:spring, 2:summer, 3:fall, 4:winter)
- yr : year (0: 2011, 1:2012)
- mnth : month ( 1 to 12)
- hr : hour (0 to 23)
- holiday : whether day is holiday or not
- weekday : day of the week (Sunday to Saturday)
- workingday : if day is neither weekend nor holiday
- weathersit :
  - 1: Clear or partly cloudy
  - 2: Mist + clouds
  - 3: Light Snow or Rain
  - 4: Heavy Rain or Snow
- temp : Normalized temperature in Celsius (divided by 41)
- atemp: Normalized feeling temperature in Celsius (divided by 50)
- hum: Normalized percent humidity (divided by 100)
- windspeed: Normalized wind speed (divided by 67)
- casual: count of casual users
- registered: count of registered users
- cnt: count of total rental bikes including casual and registered

```
In [5]: bike_trips = pd.read_csv('bikeshare.csv')

# Here we'll do some pandas datetime parsing so that the dteday column
# contains datetime objects.
bike_trips['dteday'] += ':' + bike_trips['hr'].astype(str)
bike_trips['dteday'] = pd.to_datetime(bike_trips['dteday'], format="%Y-%m-%d %H:%M")
bike_trips = bike_trips.drop(['yr', 'mnth', 'hr'], axis=1)

bike_trips.head()
```

Out[5]:

	instant	dteday	season	holiday	weekday	workingday	weathersit	temp	atemp	hum	w
0	1	2011-01-01 00:00:00	1	0	6	0	1	0.24	0.2879	0.81	
1	2	2011-01-01 01:00:00	1	0	6	0	1	0.22	0.2727	0.80	
2	3	2011-01-01 02:00:00	1	0	6	0	1	0.22	0.2727	0.80	
3	4	2011-01-01 03:00:00	1	0	6	0	1	0.24	0.2879	0.75	
4	5	2011-01-01 04:00:00	1	0	6	0	1	0.24	0.2879	0.75	



## Q1.1 (8%):

Explore the `bike_trips` dataframe to answer the following questions.

What is the data granularity? What time range is represented here? Write code in the cell below to perform your exploration.

```
In [6]: bike_trips.shape
bike_trips.iloc[0].to_frame()
bike_trips['dteday'].value_counts()
```

```
Out[6]: 2011-01-01 00:00:00    1
2012-05-03 04:00:00    1
2012-05-02 14:00:00    1
2012-05-02 15:00:00    1
2012-05-02 16:00:00    1
..
2011-09-04 01:00:00    1
2011-09-04 02:00:00    1
2011-09-04 03:00:00    1
2011-09-04 04:00:00    1
2012-12-31 23:00:00    1
Name: dteday, Length: 17379, dtype: int64
```

The granularity of bike\_trips data is that the dataframe represents the number of registered users and bike count on a certain day and hour of the year (from 2011 to 2012). This can be concluded from the dteday row which has unique values for each dteday value.

## Using pandas to plot

pandas provides useful methods on dataframes. For simple plots, we prefer to just use those methods instead of the matplotlib methods since we're often working with dataframes anyway. The syntax is:

```
dataframe.plot.<plotfunc>
```

Where the <plotfunc> is one of the functions listed under *Plotting* here:

<https://pandas.pydata.org/pandas-docs/stable/reference/frame.html>


(<https://pandas.pydata.org/pandas-docs/stable/reference/frame.html>)



In [7]: *# This plot shows the temperature at each data point*

```
bike_trips.plot.line(x='dteday', y='temp')
```

Out [7]: <AxesSubplot:xlabel='dteday'>



## seaborn

Now, we'll learn how to use the `seaborn` (<http://seaborn.pydata.org/index.html>) Python library. `seaborn` is built on top of `matplotlib` and provides many helpful functions for statistical plotting that `matplotlib` and `pandas` don't have.

Generally speaking, we'll use `seaborn` for more complex statistical plots, `pandas` for simple plots (eg. line / scatter plots), and `matplotlib` for plot customization.

Nearly all `seaborn` functions are designed to operate on `pandas` dataframes. Most of these functions assume that the dataframe is in a specific format called *long-form*, where each column of the dataframe is a particular feature and each row of the dataframe a single datapoint.

For example, this dataframe is long-form:

	country	year	avgtemp
1	Sweden	1994	6
2	Denmark	1994	6
3	Norway	1994	3
4	Sweden	1995	5
5	Denmark	1995	8
6	Norway	1995	11
7	Sweden	1996	7
8	Denmark	1996	8
9	Norway	1996	7

But this dataframe of the same data is not:

	country	avgtemp.1994	avgtemp.1995	avgtemp.1996
1	Sweden	6	5	7
2	Denmark	6	8	8
3	Norway	3	11	7


Note that the `bike_trips` dataframe is long-form.

For more about long-form data, see <https://stanford.edu/~ejdemyr/r-tutorials/wide-and-long> (<https://stanford.edu/~ejdemyr/r-tutorials/wide-and-long>). For now, just remember that we typically prefer long-form data and it makes plotting using `seaborn` easy as well.

We can create the same plot as the `dataframe.plot` above using the following `seaborn` function: `sns.lineplot`. Notice that instead of calling the plot function on dataframe object, here we pass it as a parameter to the function:

```
In [8]: # You can create a plot similar to the one above using the following s
# Note that the seaborn plot is a bit different and requires some cust
sns.lineplot(x="dteday", y="temp", data=bike_trips)
```

```
Out[8]: <AxesSubplot:xlabel='dteday', ylabel='temp'>
```



## Q1.2 (8%):

Use seaborn's `barplot` function to make a bar chart showing the average number of registered riders on each day of the week over the entire `bike_trips` dataset.

Here's a link to the seaborn API: <http://seaborn.pydata.org/api.html>  
(<http://seaborn.pydata.org/api.html>)


See if you can figure it out by reading the docs.

What trends do you notice? What do you suppose causes these trends?

Notice that `barplot` draws error bars for each category. It uses bootstrapping to make those.

```
In [9]: bike_trips["weekend"] = bike_trips["weekday"].isin([0, 6])
sns.barplot(x="weekday", y="cnt", hue="weekend", capsize=.3, data=bike
```

```
Out[9]: <AxesSubplot:xlabel='weekday', ylabel='cnt'>
```



The bar chart shows that the ride counts on weekdays are higher than on weekends. This could be caused by the fact that a lot of rented bikes are used as a transport to go to work or school. The most busy weekdays are Thursday and Friday. The error bars are the same for each category.



### Q1.3 (8%):


Now for a fancier plot that `seaborn` makes really easy to produce.

Use the `displot` function to plot a histogram of all the total rider counts in the `bike_trips` dataset. Can you fit a curve to the histogram of the data using `seaborn`? You just need to tune `kde` and `stat` parameters to the `displot` function. Fancy!

```
In [10]: # sns.displot(data=bike_trips, x="cnt", kde=True)

sns.displot(data=bike_trips, x="cnt", kde=True, stat="density", bins=100)
plt.title('Total riders count', fontsize=12)
```

```
Out[10]: Text(0.5, 1.0, 'Total riders count')
```






## Q1.4 (8%):

Use `seaborn` to make side-by-side boxplots of the number of casual riders (just checked out a bike for that day) and registered riders (have a bikeshare membership). The `boxplot` function will plot all the columns of the dataframe you pass in.

Once you make the plot, you'll notice that there are many outliers that make the plot hard to see. To mitigate this, change the y-scale to be logarithmic.

That's a plot customization so you'll use `matplotlib`. The `boxplot` function returns a `matplotlib` Axes object which represents a single plot and has a `set_yscale` function.

The result (after customization) should look like:



Revisiting Q1.1 on data granularity, what is more suitable text for `ylabel` than `Number of riders`?

```
In [11]: copy = pd.DataFrame(data={'cnt': [], 'type': []})
i = 0
for val in bike_trips["casual"]:
    copy = pd.concat([copy, pd.DataFrame([[val, 'casual']], columns=['cnt', 'type'])], axis=0)
    i += 1

for val in bike_trips["registered"]:
    copy = pd.concat([copy, pd.DataFrame([[val, 'registered']], columns=['cnt', 'type'])], axis=0)
    i += 1
copy
```


Out[11]:

	cnt	type
0	3.0	casual
1	8.0	casual
2	5.0	casual
3	3.0	casual
4	0.0	casual
...	...	...
34753	108.0	registered
34754	81.0	registered
34755	83.0	registered
34756	48.0	registered
34757	37.0	registered

34758 rows × 2 columns

```
In [12]: ax = sns.boxplot(x=copy["type"], y=copy["cnt"], order=["registered", "casual"],  
ax.set_yscale("log")  
ax.set(xlabel='', ylabel='Number of riders')
```

```
Out[12]: [Text(0.5, 0, ''), Text(0, 0.5, 'Number of riders')]
```








## Q1.5 (8%):

Let's take a closer look at the number of registered vs. casual riders.


Use the `lplot` function to make a scatterplot. Put the number of casual riders on the x-axis and the number of registered riders on the y-axis. Each point should correspond to a single row in your `bike_trips` dataframe. Notice that `seaborn` automatically fits a line of best fit to the plot. Does that line seem to be relevant?

You should note that `lplot` allows you to pass in `fit_line=False` to avoid plotting lines of best fit when you feel they are unnecessary or misleading.

There seem to be two main groups in the scatterplot. Let's see if we can separate them out. Use `lplot` to make the scatterplot again. This time, use the `hue` parameter to color points for weekday trips differently from weekend trips. You should get something that looks like:



```
In [13]: ax = sns.lmplot(x="casual", y="registered", data=bike_trips, hue='work  
legendObj = ax.legend  
handlesObj = legendObj.legendHandles  
texts = legendObj.texts  
  
plt.legend(labels=["weekend/holiday", "working day"], bbox_to_anchor=(0  
legendObj.remove())
```



## Want to learn more?

We recommend checking out the `seaborn` tutorials on your own time.

<http://seaborn.pydata.org/tutorial.html> (<http://seaborn.pydata.org/tutorial.html>)

The `matplotlib` tutorial we linked in Question 1 is also a great refresher on common

`matplotlib` functions: <https://www.labri.fr/perso/nrougier/teaching/matplotlib/>

(<https://www.labri.fr/perso/nrougier/teaching/matplotlib/>)

Here's a great blog post about the differences between Python's visualization libraries:

<https://dansaber.wordpress.com/2016/10/02/a-dramatic-tour-through-pythons-data-visualization-landscape-including-ggplot-and-altair/>


(<https://dansaber.wordpress.com/2016/10/02/a-dramatic-tour-through-pythons-data-visualization-landscape-including-ggplot-and-altair/>)

## Part 2: Self-directed EDA of social media and internet attitudes (55%)

The last part is intentionally more open-ended and will be graded on the completeness of the plot(s) produced and the insights you gain from them. The goal here is for you to thoroughly explore a dataset on attitudes towards the internet and whether it's good for society.

*Question 2.0* is asking you to look at a given visualization and reverse engineer the code that created it. *Question 2.1* is about *data exploration visualization* while the other questions are about *data presentation visualization*. Report your three most significant findings (Q2.2, Q2.3, and Q2.4). Each finding should have a *visualization headline* which highlights the main takeaway in 5-15 words, an informative visualization that supports your finding and a *visualization description*, 100-150 words per finding explaining your assumptions and what you have found. For example, the visualization headline could be "*Majority of Facebook, Instagram and Snapchat users visit these social networks daily*" with the following bar plot visualization:

Majority of Facebook, Instagram and Snapchat users visit these social networks daily



The survey data that you will analyze was collected by Pew Research. In order to access it, you need to create an account and download it from [here](https://www.pewresearch.org/internet/dataset/core-trends-survey/) (<https://www.pewresearch.org/internet/dataset/core-trends-survey/>) (click on "Download Dataset" in upper right corner). The file you will work with is January 8–February 7, 2019 – Core Trends Survey – CSV.csv. The file that contains information about the meaning of each feature is January 8–February 7, 2019 – Core Trends Survey – Questionnaire.docx.

Be sure to consider transformations, subsets, correlations, reference markers, and lines/curves-of-best-fit (as covered in Chapter 6 of PTDS) to reveal the relationship that you are wanting to learn more about. Also be sure to make plots that are appropriate for the variable types. For completeness, be explicit about any assumptions you make in your analysis. An exemplary plot will have:

- A title
- Labelled and appropriately scaled axes
- A legend, if applicable
- A carefully selected color scheme
- A main point, accentuated through design choices



## Q2.0 (5%): Reverse Engineer

Your first step is to load the data from January 8–February 7, 2019 – Core Trends Survey – CSV.csv, and understand what is stored in it. Read the the survey questionnaire in January 8–February 7, 2019 – Core Trends Survey – Questionnaire.docx to understand the meaning of each feature. Your assignment is to replicate the bar plot visualization shown above. Notice the labels on x and y axes as well as the legend of the plot to determine the information needed to construct the plot.

```
In [14]: # 1) Read your dataframe with pandas
core_trends = pd.read_csv('CoreTrendsSurvey.csv')

# 2) Identify what cols are used for plot above
core_trends_use = pd.DataFrame(core_trends[['sns2a', 'sns2b', 'sns2c',

# 3) Filter required rows and columns necessary for plotting above fig
core_trends_use.rename(columns={"sns2a": "Twitter",
                                "sns2b": "Instagram",
                                "sns2c": "Facebook",
                                "sns2d": "Snapchat",
                                "sns2e": "YouTube"}, inplace = True)

# 4) Your dataframe is in wide_format. You need to convert to long_for
# i.e. originally there is a column each for these five social medias.
core_trends_use = core_trends_use.melt()

# 5) Once dataframe is in wide format, filter individuals who responde
for column in core_trends_use:
    core_trends_use[column] = core_trends_use[column].replace(r'\s+',
    core_trends_use[column] = core_trends_use[column].replace('0', '0')
```

```

        core_trends_use[column] = core_trends_use[column].replace(9, np.nan)
        core_trends_use[column] = core_trends_use[column].replace('8', np.nan)
        core_trends_use[["value"]] = core_trends_use[["value"]].apply(pd.to_numeric)

# 6) Use pandas aggregation after groupby to calculate percentage of users
tw = core_trends_use[core_trends_use.variable == 'Twitter']
inst = core_trends_use[core_trends_use.variable == 'Instagram']
face = core_trends_use[core_trends_use.variable == 'Facebook']
snap = core_trends_use[core_trends_use.variable == 'Snapchat']
yt = core_trends_use[core_trends_use.variable == 'YouTube']

tw = tw.groupby('value', dropna=True).describe()
tw['percent'] = (tw[['variable', 'freq']] / tw[['variable', 'freq']].sum()) * 100

inst = inst.groupby('value', dropna=True).describe()
inst['percent'] = (inst[['variable', 'freq']] / inst[['variable', 'freq']].sum()) * 100

face = face.groupby('value', dropna=True).describe()
face['percent'] = (face[['variable', 'freq']] / face[['variable', 'freq']].sum()) * 100

snap = snap.groupby('value', dropna=True).describe()
snap['percent'] = (snap[['variable', 'freq']] / snap[['variable', 'freq']].sum()) * 100

yt = yt.groupby('value', dropna=True).describe()
yt['percent'] = (yt[['variable', 'freq']] / yt[['variable', 'freq']].sum()) * 100

concatenated = pd.concat([tw.assign(dataset='Twitter'),
                           inst.assign(dataset='Instagram'),
                           face.assign(dataset='Facebook'),
                           snap.assign(dataset='Snapchat'),
                           yt.assign(dataset='YouTube')
                           ])

concatenated = concatenated.reset_index()
final = pd.DataFrame(concatenated[['value', 'dataset', 'percent']])

# 7) Use seaborn barplot to plot the figure above. Customize with color
ax = sns.catplot(x = 'dataset', y = "percent", hue="value", data = final)

# 8) Add descriptive xlabel, ylabel, and title
plt.title("Majority of Facebook, Instagram and Snapchat users visit the internet several times a day")
ax.set(xlabel='Social Networks', ylabel='Percent of users per network')

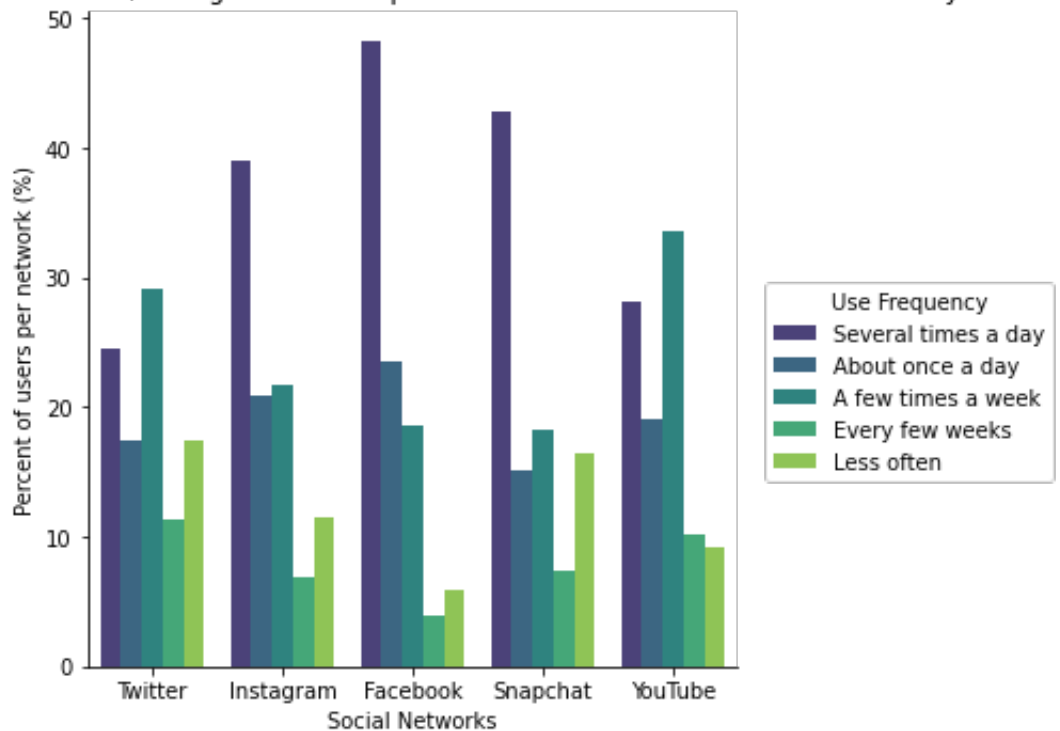
# 9) Customize legend if necessary
legendObj = ax.legend
handlesObj = legendObj.legendHandles
texts = legendObj.texts

plt.legend(labels = ['Several times a day', 'About once a day', 'A few times a week'],
           title = "Use Frequency",
           bbox to anchor=(1, 0.1, 0.5, 0.5),

```

```
handles = handlesObj)
legendObj.remove()
```

Majority of Facebook, Instagram and Snapchat users visit these social networks daily



## Q2.1 (5%): Initial exploration

Run descriptive statistics on the data by considering the EDA key data properties we covered in class. Write a 100-150 word description of your findings. Based on these statistics or other ideas you have, form hypotheses that guide your EDA and visualizations for the last three questions. You need to show at least one visualization but you are welcome to show more.

```
In [15]: # Run descriptive statistics on the data and develop ideas on what to
core_trends = pd.read_csv('CoreTrendsSurvey.csv')
core_trends.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1502 entries, 0 to 1501
Data columns (total 74 columns):
#   Column                Non-Null Count  Dtype
#   ...
```

0	respid	1502	non-null	int64
1	sample	1502	non-null	int64
2	comp	1502	non-null	int64
3	int_date	1502	non-null	int64
4	lang	1502	non-null	int64
5	cregion	1502	non-null	int64
6	state	1502	non-null	int64
7	density	1502	non-null	int64
8	sfips	1502	non-null	int64
9	usr	1502	non-null	object
10	qs1	1502	non-null	object
11	sex	1502	non-null	int64
12	eminuse	1502	non-null	int64
13	intmob	1502	non-null	int64
14	intfreq	1502	non-null	object
15	snsint2	1502	non-null	int64
16	home4nw	1502	non-null	int64
17	bbhome1	1502	non-null	object
18	bbhome2	1502	non-null	object
19	device1a	1502	non-null	object
20	smart2	1502	non-null	object
21	q20	1502	non-null	object
22	bbsmart1	1502	non-null	object
23	bbsmart2	1502	non-null	object
24	bbsmart3a	1502	non-null	object
25	bbsmart3b	1502	non-null	object
26	bbsmart3c	1502	non-null	object
27	bbsmart3d	1502	non-null	object
28	bbsmart3e	1502	non-null	object
29	bbsmart3f	1502	non-null	object
30	bbsmart3foe@	1502	non-null	object
31	bbsmart4	1502	non-null	object
32	web1a	1502	non-null	int64
33	web1b	1502	non-null	int64
34	web1c	1502	non-null	int64
35	web1d	1502	non-null	int64
36	web1e	1502	non-null	int64
37	web1f	1502	non-null	int64
38	web1g	1502	non-null	int64
39	web1h	1502	non-null	int64
40	web1i	1502	non-null	int64
41	sns2a	1502	non-null	object
42	sns2b	1502	non-null	object
43	sns2c	1502	non-null	object
44	sns2d	1502	non-null	object
45	sns2e	1502	non-null	object
46	device1b	1502	non-null	int64
47	device1c	1502	non-null	int64
48	device1d	1502	non-null	int64

```

49 books1      1502 non-null    int64
50 books2a     1502 non-null    object
51 books2b     1502 non-null    object
52 books2c     1502 non-null    object
53 age         1502 non-null    int64
54 marital     1502 non-null    int64
55 educ2       1502 non-null    int64
56 emplnw      1502 non-null    int64
57 hisp        1502 non-null    int64
58 racem1      1502 non-null    int64
59 racem2      1502 non-null    object
60 racem3      1502 non-null    object
61 racem4      1502 non-null    object
62 racecmb     1502 non-null    int64
63 birth_hisp  1502 non-null    object
64 inc         1502 non-null    int64
65 party       1502 non-null    int64
66 partyln     1502 non-null    object
67 hh1         1502 non-null    int64
68 hh3         1502 non-null    object
69 ql1         1502 non-null    object
70 ql1a        1502 non-null    object
71 qc1         1502 non-null    object
72 weight      1502 non-null    float64
73 cellweight  1502 non-null    object
dtypes: float64(1), int64(37), object(36)
memory usage: 868.5+ KB

```


```

In [16]: # Create one or more visualizations
core_trends1 = core_trends.dropna(how='any')
sns.countplot(data=core_trends1, x='educ2')
plt.show()




core_trends2 = core_trends.dropna(how='any')
sns.countplot(data=core_trends2, x='q20')
plt.show()

core_trends3 = core_trends.dropna(how='any')
sns.countplot(data=core_trends3, x='bbsmart1')
plt.show()

```







From the information above, it has been discovered that the data in the dataframe is mostly saved in either int or object formats. The int objects from educ2 describes the level of education that the participants have. This might be a good data for analyzing if level of education affects the time people use social media, or which social media is mostly used by people of certain education level. In addition to that, q20 shows on which device people usually go online. This column can be used in order to analyze the frequency of people going online with which device. The q20 column is related to BBSMART columns which in combination may show the correlation in which device users have broadband internet subscription.



## Q2.2 (15%): How education level is correlated to which social media people use

The graph below depicts the correlation of how education level affects which social media people prefer. There are 5 major social media: Twitter, Instagram, Facebook, Snapchat, and YouTube that were asked in the questionnaire. The more person uses a certain social media, the higher frequency it shows on the bar chart. The x axis shows social medias, and each column represent the education level of a person that uses it.

```
In [17]: import warnings
```

```
In [18]: warnings.filterwarnings('ignore')
edu = pd.DataFrame(core_trends[['educ2']])

core_trends_use = pd.DataFrame(core_trends[['sns2a', 'sns2b', 'sns2c',
core_trends_use.rename(columns={"sns2a": "Twitter",
                                "sns2b": "Instagram",
                                "sns2c": "Facebook",
                                "sns2d": "Snapchat",
                                "sns2e": "YouTube"}, inplace = True)

core_trends_use = core_trends_use.melt()
core_trends_use_new = core_trends_use.copy()
core_trends_use_new['educ'] = np.nan

for i in range(len(edu['educ2'])):
    val = edu['educ2'][i]
    core_trends_use_new['educ'][i + 0 * len(edu['educ2'])] = val
    core_trends_use_new['educ'][i + 1 * len(edu['educ2'])] = val
    core_trends_use_new['educ'][i + 2 * len(edu['educ2'])] = val
```

```

core_trends_use_new['educ'][i + 3 * len(edu['educ2'])] = val
core_trends_use_new['educ'][i + 4 * len(edu['educ2'])] = val

for column in core_trends_use_new:
    core_trends_use_new[column] = core_trends_use_new[column].replace(
    core_trends_use_new[column] = core_trends_use_new[column].replace(
    core_trends_use_new[column] = core_trends_use_new[column].replace(

core_trends_use_new[["educ"]] = core_trends_use_new[["educ"]].apply(pd
core_trends_use_new[["value"]] = core_trends_use_new[["value"]].apply(

for i in range(len(core_trends_use_new['value'])):
    if (core_trends_use_new['value'][i] == 1):
        core_trends_use_new['value'][i] = 5
    elif (core_trends_use_new['value'][i] == 2):
        core_trends_use_new['value'][i] = 4
    elif (core_trends_use_new['value'][i] == 4):
        core_trends_use_new['value'][i] = 2
    elif (core_trends_use_new['value'][i] == 5):
        core_trends_use_new['value'][i] = 1

ax = sns.catplot(x = 'variable', y = "value", hue="educ", data = core_
plt.gcf().set_size_inches(10, 5)

plt.title("Correlation of education level with the type of social media
ax.set(xlabel='Social Networks', ylabel='Usage frequency')

legendObj = ax.legend
handlesObj = legendObj.legendHandles
texts = legendObj.texts

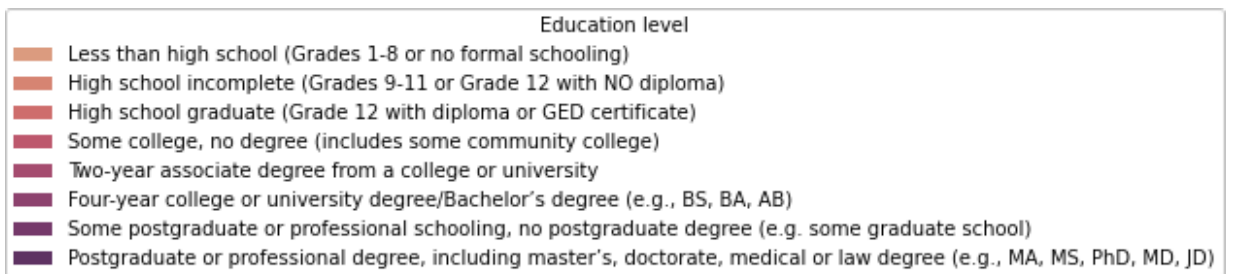
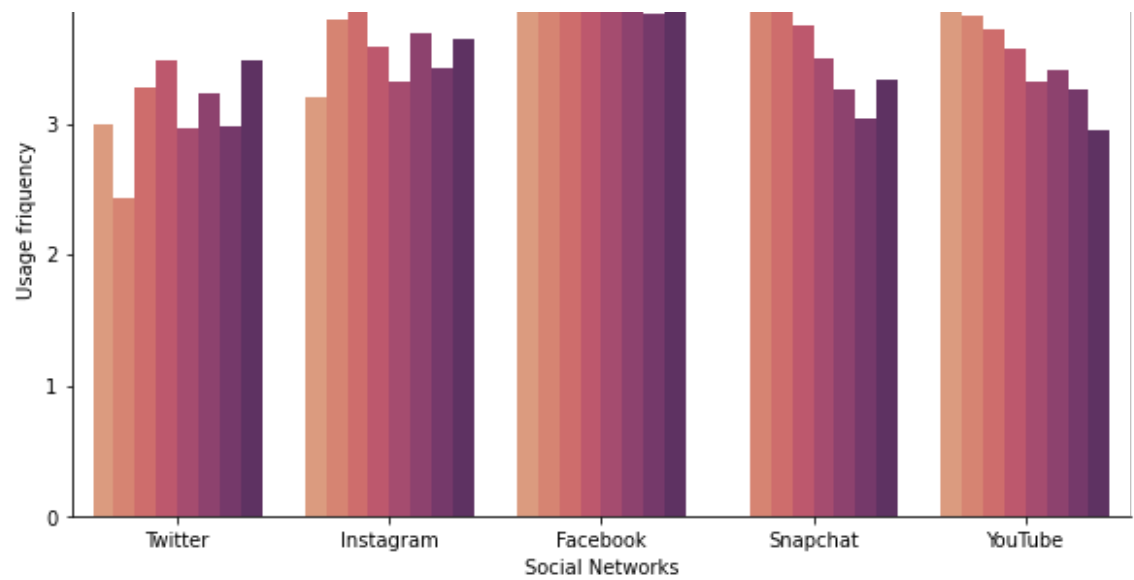
plt.legend(labels = ['Less than high school (Grades 1-8 or no formal s
                    'High school incomplete (Grades 9-11 or Grade 12
                    'High school graduate (Grade 12 with diploma or G
                    'Some college, no degree (includes some community
                    'Two-year associate degree from a college or univ
                    'Four-year college or university degree/Bachelor'
                    'Some postgraduate or professional schooling, no
                    'Postgraduate or professional degree, including m

    title = "Education level",
    loc = "lower center",
    bbox_to_anchor=(0.25, -0.6, 0.5, 0.5),
    handles = handlesObj)

legendObj.remove()

```







## **Q2.3 (15%): Smartphone owners' preferred way of accessing the internet varies substantially by education level**

The graph shows how the education level affects the type of device people usually use. The graph shows that the majority of people with a Bachelor's degree mostly use both cell phone and desktop/tablet.

Education levels:

- 1 Less than high school (Grades 1-8 or no formal schooling)
- 2 High school incomplete (Grades 9-11 or Grade 12 with NO diploma)
- 3 High school graduate (Grade 12 with diploma or GED certificate)
- 4 Some college, no degree (includes some community college)
- 5 Two-year associate degree from a college or university
- 6 Four-year college or university degree/Bachelor's degree (e.g., BS, BA, AB)
- 7 Some postgraduate or professional schooling, no postgraduate degree (e.g. some graduate school)
- 8 Postgraduate or professional degree, including master's, doctorate, medical or law degree (e.g., MA, MS, PhD, MD, JD)

```

In [19]: df = pd.read_csv('CoreTrendsSurvey.csv')

devDF = pd.DataFrame(df[['q20', 'educ2']])
pd.melt(devDF, id_vars=['q20'], value_vars=['educ2'])

for column in devDF:
    devDF[column] = devDF[column].replace(r'\s+', np.nan, regex=True)
    devDF[column] = devDF[column].replace('98', np.nan, regex=True)
    devDF[column] = devDF[column].replace('99', np.nan, regex=True)

devDF[["q20"]] = devDF[["q20"]].apply(pd.to_numeric)
devDF[["educ2"]] = devDF[["educ2"]].apply(pd.to_numeric)

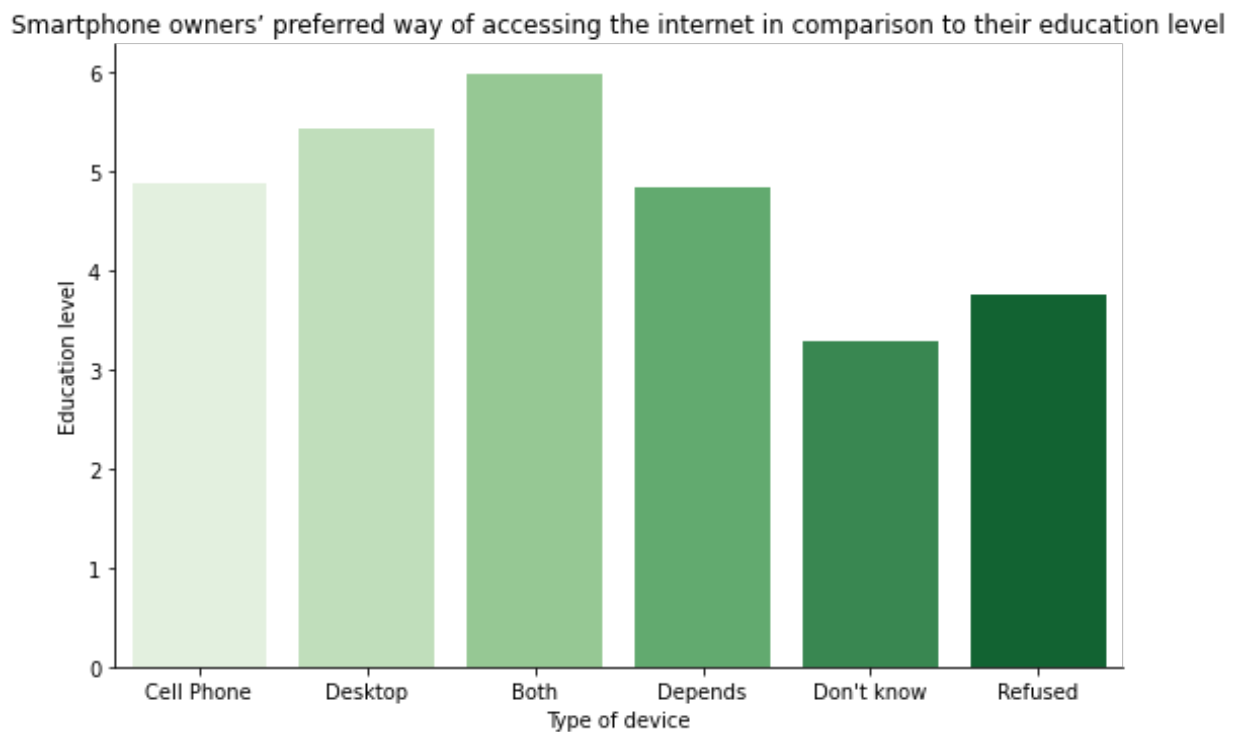
ax = sns.catplot(x = 'q20', y = "educ2", data = devDF, kind = "bar", p
plt.gcf().set_size_inches(8, 5)

ax.set_xticklabels(['Cell Phone', 'Desktop', 'Both', 'Depends', "Don't

plt.title("Smartphone owners' preferred way of accessing the internet
ax.set(xlabel='Type of device', ylabel='Education level')

```

Out[19]: <seaborn.axisgrid.FacetGrid at 0x7fa73c378fd0>





## Q2.4 (15%): How broadband subscription varies by the type of device people use

The visualization shows that whether a person had or did not have a broadband subscription, the values are the same. Probably, I did something wrong in calculations.

```
In [20]: df1 = pd.read_csv('CoreTrendsSurvey.csv')
devDF1 = pd.DataFrame(df1[['bbsmart1', 'q20']])
pd.melt(devDF1, id_vars=['bbsmart1'], value_vars=['q20'])

for column in devDF1:
    devDF1[column] = devDF1[column].replace(r'\s+', np.nan, regex=True)
    devDF1[column] = devDF1[column].replace('98', np.nan, regex=True)
    devDF1[column] = devDF1[column].replace('99', np.nan, regex=True)
    devDF1[column] = devDF1[column].replace('8', np.nan, regex=True)
    devDF1[column] = devDF1[column].replace('9', np.nan, regex=True)

devDF1[['bbsmart1']] = devDF1[['bbsmart1']].apply(pd.to_numeric)
devDF1[['q20']] = devDF1[['q20']].apply(pd.to_numeric)

ax = sns.catplot(x = 'bbsmart1', y = "q20", hue="q20", data = devDF1,
plt.gcf().set_size_inches(8, 5)

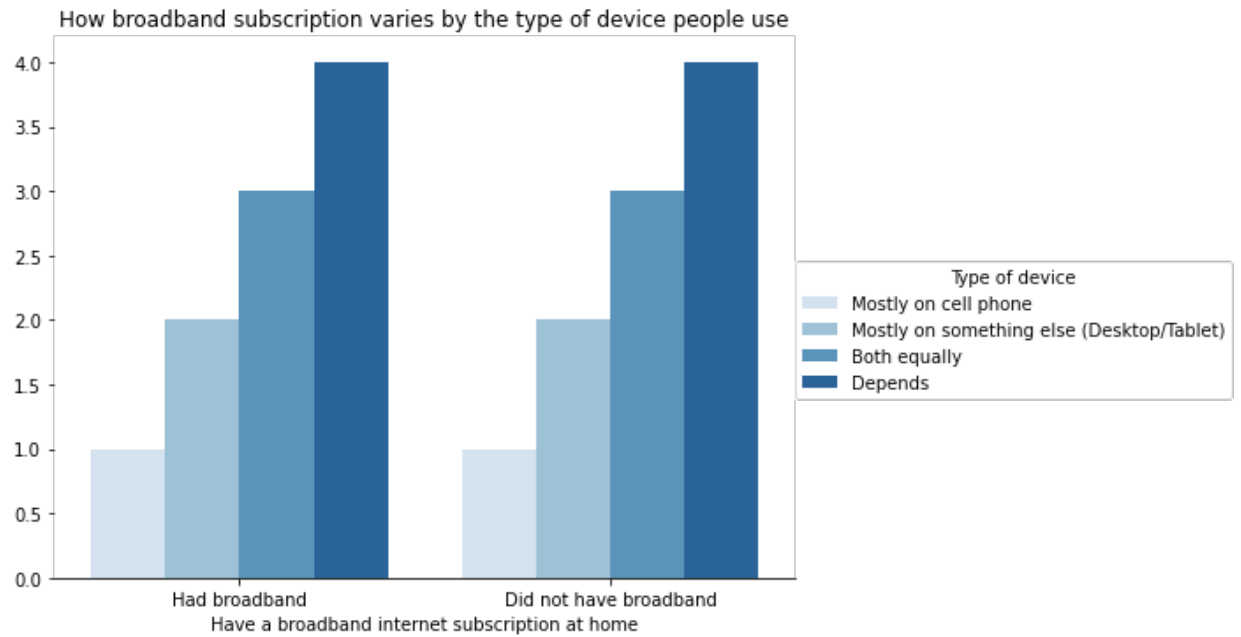
ax.set_xticklabels(['Had broadband', 'Did not have broadband'])

plt.title("How broadband subscription varies by the type of device people use")
ax.set(xlabel='Have a broadband internet subscription at home', ylabel='Broadband subscription type')
sns.despine(left=True)
sns.despine(fig=None, ax=None, top=True, right=True, left=True, bottom=True)

legendObj = ax.legend
handlesObj = legendObj.legendHandles
texts = legendObj.texts

plt.legend(labels = ['Mostly on cell phone',
                    'Mostly on something else (Desktop/Tablet)',
                    'Both equally',
                    'Depends'],
          title = "Type of device",
          loc = 1,
          bbox_to_anchor=(1.1, 0.1, 0.5, 0.5),
          handles = handlesObj)
```

```
legendObj.remove()
```



### # Extra Credit (20%)

The best 10 visualizations and insights from Questions 2.2 to 2.4 will get an extra 20% credit (at most one visualization can be considered per submission). There is nothing you need to do for the extra credit except to do your best in the last three questions. We will showcase the best visualizations in class!

This was the last part of Homework 2. Now you need to submit your work following the instructions in the beginning of the notebook and you are done!